

# MPS based on Genetic Algorithms implemented in a (Distributed) Shared-Memory Supercomputer

Oscar Peredo<sup>1</sup>, Julián M. Ortiz<sup>2</sup>

1: Computer Applications in Science & Engineering Department, Barcelona Supercomputing Center

2: Advanced Lab for Geostatistical Supercomputing, Advanced Mining Technology Center



# Outline

---

Objectives of this talk

Genetic algorithms (GAs)

Implementation

- Exploiting parallelism in GAs

- Speedup results

Conclusions

# Outline

---

## Objectives of this talk

### Genetic algorithms (GAs)

### Implementation


Exploiting parallelism in GAs

Speedup results


### Conclusions

# Objectives of this talk

- We want to generate realizations that reproduce **multiple-point statistics** (MPS) inferred from a 2D training image.



Training image




Realization A



Realization B



Realization C



Realization D

- Several methods to do this task: snesim, unilateral path, neural networks, **simulated annealing**...
- **How good are genetic algorithms to reproduce MPS?**

# Outline

---

Objectives of this talk

Genetic algorithms (GAs)

Implementation

Exploiting parallelism in GAs

Speedup results

Conclusions

# GAs: overview

---

- 1: Initial population:  $N$  random individuals
  - 2:  $k \leftarrow 0$
  - 3: Evaluate a **fitness function**  $f$  in each individual
  - 4: **while** termination criteria is not achieved **do**
  - 5:   {generation  $k$ }
  - 6:   **Selection**: select best individuals using  $f$  values
  - 7:   **Crossover**: breed new individuals crossing bits
  - 8:   **Mutation**: breed new individuals mutating some bits
  - 9:   Replace old individuals by new ones
  - 10:    $k \leftarrow k + 1$
  - 11:   Evaluate a fitness function  $f$  in each individual
  - 12: **end while**
-

# GAs: overview


---

- 1: Initial population:  $N$  random individuals
  - 2:  $k \leftarrow 0$
  - 3: Evaluate a **fitness function**  $f$  in each individual
  - 4: **while** termination criteria is not achieved **do**
  - 5:   {generation  $k$ }
  - 6:   **Selection**: select best individuals using  $f$  values
  - 7:   **Crossover**: breed new individuals crossing bits
  - 8:   **Mutation**: breed new individuals mutating some bits
  - 9:   Replace old individuals by new ones
  - 10:    $k \leftarrow k + 1$
  - 11:   Evaluate a fitness function  $f$  in each individual
  - 12: **end while**
- 

**Individuals 1D ( $nm$ )  $\longleftrightarrow$  Realization images 2D ( $n \times m$ )**

# GAs: fitness function

- It depends on a user-defined template  $T$ :



- Using  $Patterns(T)$ , we calculate:

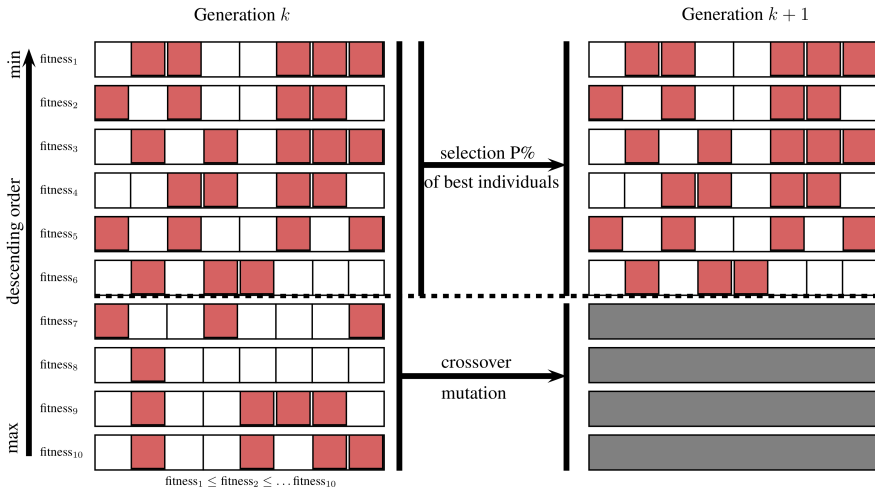
$$fitness(indiv_k) = \sum_{p \in Patterns(T)} \left( \#pattern_p^{TI} - \#pattern_p^{indiv_k} \right)^2$$

- This function measures the mismatch between the target statistics (from the training image) and the current statistics for each individual.




# GAs: selection

P% of the population "survives". The rest is generated using crossover/mutation.




# GAs: crossover


- Choose cut points



- Choose sections




- Breed new individual with the chosen sections




# GAs: mutation

- Choose mutation bit





- Breed new individual changing its value with probability  $m := m(t)$




# GAs: examples, (non)-conditional realizations


100×100 images, 10000 generations, 4000 individuals,  
20% selection, dynamic mutation probability, 100 cut-points  
Template 17 nodes disconnected




Training image




No conditionants




No conditionants




No conditionants




Training image



5% conditionants



5% conditionants




5% conditionants

# GAs: examples, dynamic template utilization

100×100 images, 10000 generations, 4000 individuals, 20% selection, dynamic mutation probability, 100 cut-points, no conditionants

Template: first half with template  $T_1$ , second half with template  $T_2$



$T_1$ : template 17 nodes



Realization gen. 5000



Realization gen. 5000




Realization gen. 5000




$T_2$ : template 9 nodes



Realization gen. 10000



Realization gen. 10000



Realization gen. 10000

# GAs: remarks

---

- Advantages:
  - Reasonable good realizations (the quality depends on the number of template's nodes).
  - straightforward and easy implementation (Fortran90)
  - the initial population can be enhanced with external information from different models to improve or refine the results.
- Disadvantages:
  - adjustment of several parameters (mutation probability, crossover cut points, population size, . . . )
  - time and memory expensive

# GAs: remarks

---

- Advantages:
  - Reasonable good realizations (the quality depends on the number of template's nodes).
  - straightforward and easy implementation (Fortran90)
  - the initial population can be enhanced with external information from different models to improve or refine the results.
- Disadvantages:
  - adjustment of several parameters (mutation probability, crossover cut points, population size, . . . )
  - **time and memory expensive** ⇒ **can we use parallel computing?**

# Outline

---

Objectives of this talk

Genetic algorithms (GAs)

Implementation

Exploiting parallelism in GAs

Speedup results

Conclusions



# Outline

---

Objectives of this talk

Genetic algorithms (GAs)

Implementation


- Exploiting parallelism in GAs

- Speedup results

Conclusions

# CPU trends

- Moore's law:  
 $2 \times \# \text{transistors every 2 years}$
- Future trends:  
more processors in a  
single chip + hardware  
accelerators




# CPU trends

---


**We have to adapt our CPU/memory intensive algorithms (like GAs) to parallel computer architectures in order to exploit all their capabilities.**

# Examples of parallel computer architectures


Memory Interconnect Cache Processor



Shared memory



Distributed memory



Distributed-Shared memory

# Programming models

---

- Shared memory: **OpenMP**  
(addition of `#pragmas`, minimal modifications in the code)
- Distributed memory: **MPI**  
(synchronization of low-level network messages between nodes, considerable modifications in the code)
- Distributed-shared memory: **MPI+OpenMP**

# Programming models

---

- Shared memory: **OpenMP**  
(addition of `#pragmas`, minimal modifications in the code)
- Distributed memory: **MPI**  
(synchronization of low-level network messages between nodes, considerable modifications in the code)
- Distributed-shared memory: **MPI+OpenMP** (best performance)

# Exploiting parallelism in GAs

---

**Task:** we need to calculate the fitness function of  $N$  individuals.

**Elements:**  $P$  nodes,  $C$  processors per node.

**Proposed solution:** MPI+OpenMP based solution

- Distribute the population in  $N/P$  individuals per node (MPI).
- Each node calculates the fitness function of  $N/P$  individuals using  $C$  processors (OpenMP parallel loops).
- Each node shares its best individuals with the other nodes (MPI).

# Outline

---

Objectives of this talk

Genetic algorithms (GAs)

Implementation

Exploiting parallelism in GAs

Speedup results

Conclusions



# Speedup results

- **Hardware:** 126 nodes (12 CPUs, 2.53 GHz), 24 GB RAM/node.
- **OpenMP:** fitness function for an individual of size  $1000 \times 1000$ , 17-node template

CPUs	Time(secs)	Speedup
1	0.402347	1x
1+code opts	0.243200	1.65x
2+code opts	0.1280000	3.14x
4+code opts	0.076800004	5.23x
6+code opts	0.044799998	8.98x
8+code opts	0.038400002	10.47x
10+code opts	0.032000002	12.57x
12+code opts	0.025599999	15.71x

\*code opts: techniques adapted from stencil optimization in Finite Differences methods, like memory accesses minimization, SIMDization, blocking, . . .

# Speedup results

- **Hardware:** 126 nodes (12 CPUs, 2.53 GHz), 24 GB RAM/node.
- **MPI+OpenMP:** 30 generations, 1000 individuals of size  $1000 \times 1000$ , 17-node template

Nodes/CPU	Time(secs)	Speedup
1/1	9364	1x
1/12	1525	6.14x
2/12	781	11.98x
4/12	394	23.76x
8/12	203	46.12x

# Outline

---

Objectives of this talk

Genetic algorithms (GAs)

Implementation

Exploiting parallelism in GAs

Speedup results

Conclusions

# Conclusions

---

- Easy implementation and parallelization: unfortunately it takes too much time to converge...
- Dynamic mutation rate (annealing-style) and dynamic template accelerates the convergence of MPS-GAs.
- Refinement of realizations: adding realizations generated with other methods to the initial population.
- *Future work*: inclusion of hardware accelerators (GPUs) with new pragma-based programming models

MPI+OpenMP+{OpenACC,OmpSs,...}

Thanks for your attention!

**Contact:** oscar.peredo [at] bsc.es

**Contact:** jortiz [at] ing.uchile.cl

<http://www.bsc.es>

<http://www.alges.cl>

<http://www.amtc.cl>