

# Resurrecting GSLIB by code optimization and multi-core programming

Oscar Peredo<sup>1</sup> & Julián M. Ortiz<sup>1,2</sup>

<sup>1</sup>Advanced Lab for Geostatistical Supercomputing, Advanced Mining Technology Center, University of Chile

<sup>2</sup>Department of Mining Engineering, University of Chile



16th Annual Conference of the International Association for Mathematical Geosciences,

17-20 October, 2014 - New Delhi, India

# Outline

---

## Motivation

## Methodology

Step 1: Re-design

Step 2: Profiling and code optimization

Step 3: Multi-core execution

## Conclusions and Future work

# Outline

---

## Motivation

## Methodology

Step 1: Re-design

Step 2: Profiling and code optimization

Step 3: Multi-core execution

## Conclusions and Future work

# Motivation

---

- GSLIB: open-source software package (*applications + utilities*) used in the geostatistical community for more than 30 years.

# Motivation

---

- GSLIB: open-source software package (*applications + utilities*) used in the geostatistical community for more than 30 years.
- Significant efforts have been developed to accelerate/enhance the scope of the original package (SGEMS, WinGSLIB), but according to the authors' knowledge, few efforts have been reported in order to accelerate the GSLIB package by itself.

# Motivation

---

- GSLIB: open-source software package (*applications + utilities*) used in the geostatistical community for more than 30 years.
- Significant efforts have been developed to accelerate/enhance the scope of the original package (SGEMS, WinGSLIB), but according to the authors' knowledge, few efforts have been reported in order to accelerate the GSLIB package by itself.
- How can we accelerate an application *by itself*?  
***Optimizing its code and using multiple threads (cores) of execution***

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:



# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*
  - Inlining

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*
  - Inlining
  - Reduction of miss-predicted braches

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*
  - Inlining
  - Reduction of miss-predicted braches
  - Loop unrolling

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*
  - Inlining
  - Reduction of miss-predicted braches
  - Loop unrolling
  - Data locality and memory bandwidth

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*
  - Inlining
  - Reduction of miss-predicted braches
  - Loop unrolling
  - Data locality and memory bandwidth
  - SIMD programming

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*
  - Inlining
  - Reduction of miss-predicted braches
  - Loop unrolling
  - Data locality and memory bandwidth
  - SIMD programming
  - ...

# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*
  - Inlining
  - Reduction of miss-predicted braches
  - Loop unrolling
  - Data locality and memory bandwidth
  - SIMD programming
  - ...



# What is *optimize code*?

---

- *Code optimization*: set of techniques and code modifications (machine-independent and machine-dependent) that can help us to accelerate an application using a general-purpose CPU.
- Examples:
  - *Memoization*
  - Inlining
  - Reduction of miss-predicted braches
  - Loop unrolling
  - Data locality and memory bandwidth
  - SIMD programming
  - ...

GSLIB + Code optimization

=

**faster GSLIB!**

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:
  - OpenMP

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:
  - OpenMP
  - POSIX threads

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:
  - OpenMP
  - POSIX threads
  - MPI (intra-node)

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:
  - OpenMP
  - POSIX threads
  - MPI (intra-node)
  - Intel's TBB

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:
  - OpenMP
  - POSIX threads
  - MPI (intra-node)
  - Intel's TBB
  - Erlang (intra-node)



# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:
  - OpenMP
  - POSIX threads
  - MPI (intra-node)
  - Intel's TBB
  - Erlang (intra-node)
  - ...

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:
  - OpenMP
  - POSIX threads
  - MPI (intra-node)
  - Intel's TBB
  - Erlang (intra-node)
  - ...

# What is *multi-core programming*?

---

- *Multi-core programming*: set of techniques and programming models that can help us to accelerate an application using multiple general-purpose CPUs, typically sharing the same memory address space (intra-node communication).
- Examples:
  - OpenMP
  - POSIX threads
  - MPI (intra-node)
  - Intel's TBB
  - Erlang (intra-node)
  - ...

GSLIB + Code optimization + Multi-core programming

=

**even faster GSLIB!!!**

# Objective

---

Develop a new open-source version of GSLIB, that uses optimized code and multiple threads (cores) of execution.

# Outline

---

## Motivation

## Methodology

Step 1: Re-design

Step 2: Profiling and code optimization

Step 3: Multi-core execution

## Conclusions and Future work

# Methodology

---

1. Pick an application and re-design it for acceleration.

# Methodology

---

1. Pick an application and re-design it for acceleration.
2. Create a run-time profile, identifying bottle-necks and optimize those bottle-necks.

# Methodology

---

1. Pick an application and re-design it for acceleration.
2. Create a run-time profile, identifying bottle-necks and optimize those bottle-necks.
3. Adapt the application for multi-thread execution.



# Methodology

---

1. Pick an application and re-design it for acceleration.
2. Create a run-time profile, identifying bottle-necks and optimize those bottle-necks.
3. Adapt the application for multi-thread execution.

# Methodology

---

1. Pick an application and re-design it for acceleration.
2. Create a run-time profile, identifying bottle-necks and optimize those bottle-necks.
3. Adapt the application for multi-thread execution.

Let's see these steps using `gamv` as example...

# Outline

---

Motivation

Methodology

Step 1: Re-design

Step 2: Profiling and code optimization

Step 3: Multi-core execution

Conclusions and Future work

# Step 1: Re-design

- Inline important subroutines (readparm and gamv in this case):

```

subroutine foo(a)
  integer a
  a=a+2
end subroutine foo

```

```

program main
  integer a,b
  a=1
  call foo(a)
  b=a
end program main

```

```

csubroutine foo(a)
c  integer a
c  a=a+2
cend subroutine foo

```

```

program main
  integer a,b
  a=1
c  call foo(a)
  b=a+2
end program main

```

# Step 1: Re-design

- Inline important subroutines (readparm and gamv in this case):

```
subroutine foo(a)
  integer a
  a=a+2
end subroutine foo
```

```
program main
  integer a,b
  a=1
  call foo(a)
  b=a
end program main
```

```
csubroutine foo(a)
c  integer a
c  a=a+2
cend subroutine foo
```

```
program main
  integer a,b
  a=1
c  call foo(a)
  b=a+2
end program main
```

- Move global variables from module geostat and common blocks (include file.inc) to main program.

# Step 1: Re-design

- Inline important subroutines (readparm and gamv in this case):

```
subroutine foo(a)
  integer a
  a=a+2
end subroutine foo
```

```
program main
  integer a,b
  a=1
  call foo(a)
  b=a
end program main
```

```
csubroutine foo(a)
c  integer a
c  a=a+2
cend subroutine foo
```

```
program main
  integer a,b
  a=1
c  call foo(a)
  b=a+2
end program main
```

- Move global variables from module geostat and common blocks (include file.inc) to main program.
- Use implicit none in all non-inlined subroutines.

# Step 1: Re-design

- Inline important subroutines (`readparm` and `gamv` in this case):

```
subroutine foo(a)
  integer a
  a=a+2
end subroutine foo
```

```
program main
  integer a,b
  a=1
  call foo(a)
  b=a
end program main
```

```
csubroutine foo(a)
c  integer a
c  a=a+2
cend subroutine foo
```

```
program main
  integer a,b
  a=1
c  call foo(a)
  b=a+2
end program main
```

- Move global variables from module `geostat` and common blocks (`include file.inc`) to main program.
- Use `implicit none` in all non-inlined subroutines.
- Avoid side-effects in each subroutine, passing all necessary data as explicit subroutine parameters.

# Step 1: Re-design

- Inline important subroutines (`readparm` and `gamv` in this case):

```
subroutine foo(a)
  integer a
  a=a+2
end subroutine foo
```

```
program main
  integer a,b
  a=1
  call foo(a)
  b=a
end program main
```

```
csubroutine foo(a)
c  integer a
c  a=a+2
cend subroutine foo
```

```
program main
  integer a,b
  a=1
c  call foo(a)
  b=a+2
end program main
```

- Move global variables from module `geostat` and common blocks (`include file.inc`) to main program.
- Use `implicit none` in all non-inlined subroutines.
- Avoid side-effects in each subroutine, passing all necessary data as explicit subroutine parameters.



# Step 1: Re-design

- Inline important subroutines (`readparm` and `gamv` in this case):

```
subroutine foo(a)
  integer a
  a=a+2
end subroutine foo
```

```
program main
  integer a,b
  a=1
  call foo(a)
  b=a
end program main
```

```
csubroutine foo(a)
c  integer a
c  a=a+2
cend subroutine foo
```

```
program main
  integer a,b
  a=1
c  call foo(a)
  b=a+2
end program main
```

- Move global variables from module `geostat` and common blocks (`include file.inc`) to main program.
- Use `implicit none` in all non-inlined subroutines.
- Avoid side-effects in each subroutine, passing all necessary data as explicit subroutine parameters.

Maybe the re-designed code doesn't look pretty, but the ***data-flow path*** of each variable can be easily identified (important for multi-core programming).

## Example: gamv

Original code:	Re-designed code:
<pre> module geostat   real, allocatable::     x(:),y(:),z(:)   ... end module geostat program main   use geostat   call readparm   call gamv   call writeout   stop end end program main </pre>	<pre> program main <b>c move module vars into main</b>   real, allocatable::     x(:),y(:),z(:)   ... <b>c inlined code of readparm</b>   ... <b>c inlined code of gamv</b>   ... <b>c writeout is not inlined</b>   call writeout   stop end end program main </pre>

# Outline

---

Motivation

## Methodology

Step 1: Re-design

Step 2: Profiling and code optimization

Step 3: Multi-core execution

Conclusions and Future work

## Step 2: Profiling and code optimization

---

- Compile the code using the *debug* flag ( $-g$ ).

## Step 2: Profiling and code optimization

---

- Compile the code using the *debug* flag (`-g`).
- Select a Fortran profiler tool (well-known tools for Linux: `gprof` and `oprofile`).

## Step 2: Profiling and code optimization

---

- Compile the code using the *debug* flag (`-g`).
- Select a Fortran profiler tool (well-known tools for Linux: `gprof` and `oprofile`).
- Identify bottle-necks looking at the performance profiles obtained (sampling in hardware events).

## Step 2: Profiling and code optimization

---

- Compile the code using the *debug* flag (`-g`).
- Select a Fortran profiler tool (well-known tools for Linux: `gprof` and `oprofile`).
- Identify bottle-necks looking at the performance profiles obtained (sampling in hardware events).
- Optimize the code removing or minimizing those bottle-necks.

Event=BR\_MISS\_PRED\_RETIRED, sampling freq. 500 hw events.

15/29



Event=BR\_MISS\_PRED\_RETIRED, sampling freq. 500 hw events.

```

: ...
: if (hs.gt.dismxs) go to 4
1663 14.7954 : if (hs.lt.0.0) hs = 0.0
1 0.0089 : h = sqrt (hs)
: c
: c Determine which lag this is and skip if outside the defined
: c distance tolerance:
: c
: if (h.le.EPSLON) then
: lagbeg = 1
1 0.0089 : lagend = 1
: else
: lagbeg = -1
: lagend = -1
2510 22.3310 : do ilag=2,nlag+2
3813 33.9235 : if (h.ge. (xlag*real (ilag-2)-xltol) .and.
: h.le. (xlag*real (ilag-2)+xltol)) then
: +
254 2.2598 : if (lagbeg.lt.0) lagbeg = ilag
4 0.0356 : lagend = ilag
: end if
: end do
1 0.0089 : if (lagend.lt.0) go to 4
: endif

```

Too many "if" instructions

# Too many "if" instructions

Event=BR\_MISS\_PRED\_RETIRED, sampling freq. 500 hw events.

```

: ...
:         if (hs.gt.dismxs) go to 4
1452 17.8356 :         if (hs.lt.0.0) hs = 0.0
      2 0.0246 :         h = sqrt(hs)
:
: c Determine which lag this is and skip if outside the defined
: c distance tolerance:
: c
:         if (h.le.EPSLON) then
:             lagbeg = 1
:             lagend = 1
:         else
:             lagbeg = -1
:             lagend = -1
:             liminf=(ceiling((h-xltol)*xlaginv)+2)
270 3.3165 :             limsup=(floor((h+xltol)*xlaginv)+2)
1022 12.5537 :             do ilag=liminf,limsup
      700 8.5985 :                 if (lagbeg.lt.0) lagbeg = ilag
      350 4.2992 :                 lagend = ilag
:
:             end do
350 4.2992 :             if (lagend.lt.0) go to 4
:         endif
:

```

# Example: oprofile output over gamv optimized

Event=BR\_MISS\_PRED\_RETIRED, sampling freq. 500 hw events.

#events	%total	line of code
-----		
		...
		if(hs.gt.dismxs) go to 4
1452	17.8356	if(hs.lt.0.0) hs = 0.0
2	0.0246	h = sqrt(hs)
		:c
		:c Determine which lag this is and skip if outside the defined
		:c distance tolerance:
		:c
		if(h.le.EPSILON) then
		lagbeg = 1
		lagend = 1
		else
		lagbeg = -1
		lagend = -1
		liminf=(ceiling((h-xltol)*xlagnv)+2)
270	3.3165	limsup=(floor((h+xltol)*xlagnv)+2)
1022	12.5537	do ilag=liminf,limsup
700	8.5985	if(lagbeg.lt.0) lagbeg = ilag
350	4.2992	lagend = ilag
		:
		:
		end do
350	4.2992	if(lagend.lt.0) go to 4
		endif
		...

Acceleration  
6%

# Example: oprofile output over gamv

Event=L1D\_BLOCKS, sampling freq. 100000 hw events.

```
#events %total : line of code
```

```
-----
      :...
      :c MAIN LOOP OVER ALL PAIRS:
      :c
      :      irepo = max(1,min((nd/10),1000))
      :      do 3 i=1,nd
      :          if((int(i/irepo)*irepo).eq.i) write(*,103) i,nd
      :
      :
      :
488  2.7767 :      do 4 j=i,nd
      :c
      :c Definition of the lag corresponding to the current pair:
      :c
1500  8.5349 :          dx = x(j) - x(i)
901   5.1266 :          dy = y(j) - y(i)
1666  9.4794 :          dz = z(j) - z(i)
88   0.5007 :          dxs = dx*dx
811  4.6145 :          dys = dy*dy
1964 11.1750 :          dzs = dz*dz
5610 31.9203 :          hs = dxs + dys + dzs
3837 21.8321 :          if(hs.gt.dismxs) go to 4
8    0.0455 :          if(hs.lt.0.0) hs = 0.0
2    0.0114 :          h = sqrt(hs)
      :...

```

# Example: oprofile output over gamv

Event=L1D\_BLOCKS, sampling freq. 100000 hw events.

```
#events %total : line of code
```

```
-----
: ...
:c MAIN LOOP OVER ALL PAIRS:
:c
:       irepo = max(1,min((nd/10),1000))
:       do 3 i=1,nd
:           if((int(i/irepo)*irepo).eq.i) write(*,103) i,nd
:
:
:
488  2.7767 :       do 4 j=i,nd
:c
:c Definition of the lag corresponding to the current pair:
:c
1500  8.5349 :       dx = x(j) - x(i)
901   5.1266 :       dy = y(j) - y(i)
1666  9.4794 :       dz = z(j) - z(i)
88   0.5007 :       dxs = dx*dx
811  4.6145 :       dys = dy*dy
1964 11.1750 :       dzs = dz*dz
5610 31.9203 :       hs = dxs + dys + dzs
3837 21.8321 :       if(hs.gt.dismxs) go to 4
8    0.0455 :       if(hs.lt.0.0) hs = 0.0
2    0.0114 :       h  = sqrt(hs)
: ...
```

Too many  
memory  
accesses

# Example: oprofile output over gamv optimized

Event=L1D\_BLOCKS, sampling freq. 100000 hw events.

```
#events %total : line of code
-----
: ...
:c MAIN LOOP OVER ALL PAIRS:
:c
:      irepo = max(1,min((nd/10),1000))
:      do 3 i=1,nd
:          if((int(i/irepo)*irepo).eq.i) write(*,103) i,nd
:              xi=x(i)
:              yi=y(i)
:              zi=z(i)
291  1.3619 :      do 4 j=i,nd
:c
:c Definition of the lag corresponding to the current pair:
:c
802  3.7535 :          dx = x(j) - xi
1040 4.8673 :          dy = y(j) - yi
1368 6.4024 :          dz = z(j) - zi
785  3.6739 :          dxs = dx*dx
595  2.7847 :          dys = dy*dy
148  0.6927 :          dzs = dz*dz
5881 27.5238 :          hs = dxs + dys + dzs
3574 16.7267 :          if(hs.gt.dismxs) go to 4
6  0.0281 :          if(hs.lt.0.0) hs = 0.0
:          h = sqrt(hs)
: ...
```

# Example: oprofile output over gamv optimized

Event=L1D\_BLOCKS, sampling freq. 100000 hw events.

```
#events %total : line of code
```

```
-----
: ...
:c MAIN LOOP OVER ALL PAIRS:
:c
:       irepo = max(1,min((nd/10),1000))
:       do 3 i=1,nd
:           if((int(i/irepo)*irepo).eq.i) write(*,103) i,nd
:           xi=x(i)
:           yi=y(i)
:           zi=z(i)
291  1.3619 :       do 4 j=i,nd
:c
:c Definition of the lag corresponding to the current pair:
:c
802  3.7535 :           dx = x(j) - xi
1040 4.8673 :           dy = y(j) - yi
1368 6.4024 :           dz = z(j) - zi
785  3.6739 :           dxs = dx*dx
595  2.7847 :           dys = dy*dy
148  0.6927 :           dzs = dz*dz
5881 27.5238 :           hs = dxs + dys + dzs
3574 16.7267 :           if(hs.gt.dismxs) go to 4
6  0.0281 :           if(hs.lt.0.0) hs = 0.0
:           h = sqrt(hs)
: ...
```

**Acceleration  
22%**

# Outline

---

Motivation

## Methodology

Step 1: Re-design

Step 2: Profiling and code optimization

Step 3: Multi-core execution

Conclusions and Future work



## Step 3: Multi-core execution

---

- We choose the OpenMP programming model: minimal modifications in the original code and easy implementation.

## Step 3: Multi-core execution

---

- We choose the OpenMP programming model: minimal modifications in the original code and easy implementation.
- Compile the code using the openmp flag (`-fopenmp` for GCC or `-openmp` for Intel compilers).

## Step 3: Multi-core execution

---

- We choose the OpenMP programming model: minimal modifications in the original code and easy implementation.
- Compile the code using the openmp flag (`-fopenmp` for GCC or `-openmp` for Intel compilers).
- Identify the most time consuming parts of the code, looking the previous performance profiles obtained.

## Step 3: Multi-core execution

---

- We choose the OpenMP programming model: minimal modifications in the original code and easy implementation.
- Compile the code using the openmp flag (`-fopenmp` for GCC or `-openmp` for Intel compilers).
- Identify the most time consuming parts of the code, looking the previous performance profiles obtained.
- Add OpenMP directives in order to parallelize those parts using multiple threads of execution (ideally, 1 thread must run in 1 CPU core).

# Example: +100% faster gamv

```

...
c
c MAIN LOOP OVER ALL PAIRS:
c
      irepo = max(1,min((nd/10),1000))

      do 3 i=1,nd
        if((int(i/irepo)*irepo).eq.i) write(*,103) i,nd
          xi=x(i)
          yi=y(i)
          zi=z(i)
      do 4 j=i,nd
c
c Definition of the lag corresponding to the current pair:
c
          dx  = x(j) - xi
          dy  = y(j) - yi
          dz  = z(j) - zi
          dxs = dx*dx
          dys = dy*dy
          dzs = dz*dz
      ...
...

```

# Example: +100% faster gamv

```

c
c MAIN LOOP OVER ALL PAIRS:
c
      irepo = max(1,min((nd/10),1000))

do 3 i=1,nd
  if((int(i/irepo)*irepo).eq.i) write(*,103) i,nd
  xi=x(i)
  yi=y(i)
  zi=z(i)
do 4 j=i,nd
c
c Definition of the lag corresponding to the current pair:
c
      dx = x(j) - xi
      dy = y(j) - yi
      dz = z(j) - zi
      dxs = dx*dx
      dys = dy*dy
      dzs = dz*dz
...

```

$$\frac{nd*(nd+1)}{2}$$

pairs of points  
to process

# Example: +100% faster gamv + OpenMP

```

...
c
c MAIN LOOP OVER ALL PAIRS:
c
      irepo = max(1,min((nd/10),1000))
c$omp parallel default(firstprivate) shared(x,y,z,reducedVariables)
#ifdef _OPENMP
      threadId = int(OMP_get_thread_num())+1
#endif
c$omp do schedule(runtime)
      do 3 i=1,nd
        if((int(i/irepo)*irepo).eq.i) write(*,103) i,nd
        xi=x(i)
        yi=y(i)
        zi=z(i)
      do 4 j=1,nd
c
c Definition of the lag corresponding to the current pair:
c
        dx = x(j) - xi
        dy = y(j) - yi
        dz = z(j) - zi
        dxs = dx*dx
        dys = dy*dy
        dzs = dz*dz
      ...
c$ omp end do
c$ omp end parallel
...

```

# Example: +100% faster gamv + OpenMP

```

c
c MAIN LOOP OVER ALL PAIRS:
c
      irepo = max(1,min((nd/10),1000))
c$omp parallel default(firstprivate) shared(x,y,z,reducedVariables)
#ifdef _OPENMP
      threadId = int(OMP_get_thread_num())+1
#endif
c$omp do schedule(runtime)
      do 3 i=1,nd
        if((int(i/irepo)*irepo).eq.i) write(*,103) i,nd
        xi=x(i)
        yi=y(i)
        zi=z(i)
        do 4 j=i,nd
c
c Definition of the lag corresponding to the current pair:
c
          dx = x(j) - xi
          dy = y(j) - yi
          dz = z(j) - zi
          dxs = dx*dx
          dys = dy*dy
          dzs = dz*dz
          ...
c$ omp end do
c$ omp end parallel

```

Outer-loop is divided among  $T$  threads using a load-balance policy defined at runtime



# Example: +100% faster gamv + OpenMP

Runtime policy: "static" (default)

Each thread process  $nd/T$  contiguous iterations

$T = 2$




$T = 4$



$T = 8$



$T = 16$



Can we run faster?

# Example: +100% faster gamv + OpenMP

Runtime policy: "static" (default)

Each thread process  $nd/T$  contiguous iterations

$T = 2$




$T = 4$



$T = 8$



$T = 16$



Can we run faster?

# Example: +100% faster gamv + OpenMP

Runtime policy: "static,16"

Each thread process 16 contiguous iterations (interleaved)

$T = 2$



$T = 4$



$T = 8$




$T = 16$



Yes!

# Other results: gamv, sisim and sgsim



# Outline

---

## Motivation

## Methodology

Step 1: Re-design

Step 2: Profiling and code optimization

Step 3: Multi-core execution

## Conclusions and Future work

# Conclusions

---

- **The proposed methodology accelerates the runtime execution of GSLIB applications, without changing the classical usability (parameter files).**
- Many code optimization techniques can be applied to the target code, in a combined way.
- Different speedup results can be obtained using different load-balancing policies, so we must choose the best for our problem.
- Current developments: `gamv`, `sisim` and `sgsim`
  - Linux: openSUSE 12.2, GNU Fortran compiler 4.7 (`gfortran`)
  - Windows: Windows 7, Visual Studio 2013, Intel Parallel Studio XE 2013 (`ifort.exe`)

# Future work

---

- Apply this methodology to the most important GSLIB applications and utilities.
- Other GSLIB-based applications can also be benefited from this effort (following a similar acceleration methodology).
- If more processing power is needed, we can use **hardware acceleration devices** or **distributed computing**:
  - Early experiments using **Intel Xeon Phi (MICs)** co-processors (60 cores of execution).
  - **Nvidia CUDA-based GSLIB** which can use GPUs as back-end compute engines.
  - For really large computations, we are planning to implement an **MPI-based GSLIB** which can run in a cluster of distributed general-purpose computers.

# Thanks for your attention!

Check for new updates and the latest release in:

`http://gslib.alges.cl`

**Contact:** `operedo [at] alges.cl`