

Procesamiento masivo de datos con Spark

Oscar Peredo A.
Telefónica I+D Chile
oscar.peredo@telefonica.com

2 Junio 2016



ELEGIMOS TODO_

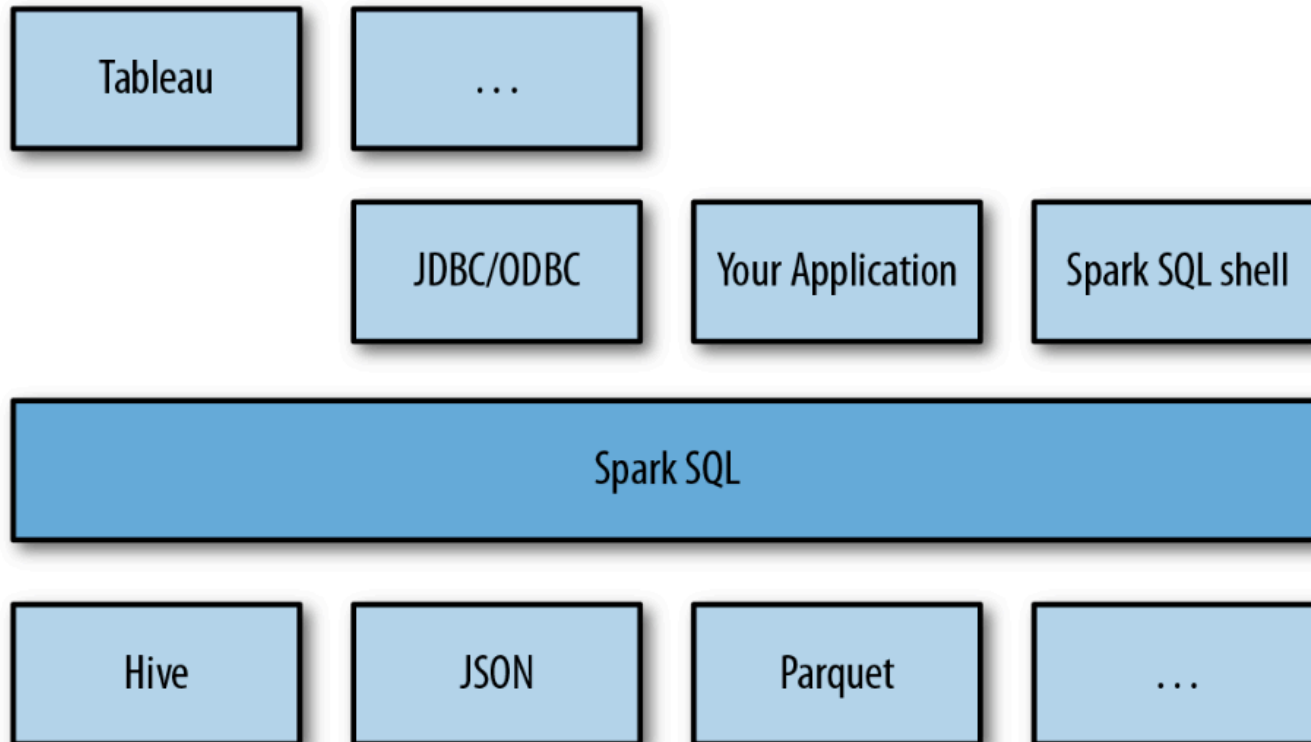
Contenidos

- Spark SQL
- Tez versus Spark

Spark SQL

Herramienta útil cuando:

- Los datos (semi) estructurados
- La base de conocimiento en SQL es amplia
- Se necesita integración con equipo de data scientists (Python)
- Se quiere usar toda la potencia de procesamiento de Spark como backend de cómputo



Ejemplo: SparkSQL + Hive

Por el lado de Hive (datos)...

```
profesor@hn0-ejempl:~$ hive
WARNING: Use "yarn jar" to launch YARN applications.

Logging initialized using configuration in file:/etc/hive/2.4.2.0-258/0/hive-log4j.properties
hive> show tables from default;
OK
hivesampletable
Time taken: 2.974 seconds, Fetched: 1 row(s)
```

Ejemplo: SparkSQL + Hive

```
hive> show columns in hivesampletable;  
OK  
clientid  
querytime  
market  
deviceplatform  
devicemake  
devicemodel  
state  
country  
querydwelltime  
sessionid  
sessionpagevieworder  
Time taken: 0.447 seconds, Fetched: 11 row(s)  
hive> █
```

Ejemplo: SparkSQL + Hive

```
hive> show tables from default;
OK
hivesampletable
Time taken: 3.331 seconds, Fetched: 1 row(s)
hive> select * from hivesampletable limit 10;
OK
8      18:54:20      en-US  Android Samsung SCH-i500      California      United States      13.9204007      0      0
23     19:19:44      en-US  Android HTC      Incredible      Pennsylvania      United States      NULL      0      0
23     19:19:46      en-US  Android HTC      Incredible      Pennsylvania      United States      1.4757422      0      1
23     19:19:47      en-US  Android HTC      Incredible      Pennsylvania      United States      0.245968      0      2
28     01:37:50      en-US  Android Motorola      Droid X Colorado      United States      20.3095339      1      1
28     00:53:31      en-US  Android Motorola      Droid X Colorado      United States      16.2981668      0      0
28     00:53:50      en-US  Android Motorola      Droid X Colorado      United States      1.7715228      0      1
28     16:44:21      en-US  Android Motorola      Droid X Utah      United States      11.6755987      2      1
28     16:43:41      en-US  Android Motorola      Droid X Utah      United States      36.9446892      2      0
28     01:37:19      en-US  Android Motorola      Droid X Colorado      United States      28.9811416      1      0
Time taken: 1.991 seconds, Fetched: 10 row(s)
```

Ejemplo: SparkSQL + Hive

Por el lado de SparkSQL (query)...

```
from pyspark import SparkConf, SparkContext, SQLContext
from pyspark.sql.types import StringType

# Import Spark SQL
from pyspark.sql import HiveContext, Row
# Or if you can't include the hive requirements
from pyspark.sql import SQLContext, Row

conf = SparkConf().setAppName("GenerateTraces")
sc = SparkContext(conf = conf)
hiveCtx = HiveContext(sc)
query = hiveCtx.sql("SELECT count(*) FROM hivesampletable")
for x in query.collect():
    print x
```


Ejemplo: SparkSQL + Hive

```
profesor@hn0-ejempl:~$ spark-submit querySparkSQL.py
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hdp/2.4.2.0-258/spark/lib/spark-assembly-1.6.1.2.4.2.0-258.jar:]
SLF4J: Found binding in [jar:file:/usr/hdp/2.4.2.0-258/spark/lib/spark-examples-1.6.1.2.4.2.0-258.jar:]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
16/06/02 21:32:39 INFO SparkContext: Running Spark version 1.6.1
16/06/02 21:32:40 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform.
16/06/02 21:32:40 INFO SecurityManager: Changing view acls to: profesor
16/06/02 21:32:40 INFO SecurityManager: Changing modify acls to: profesor
...
16/06/02 21:34:18 INFO YarnScheduler: Removed TaskSet 1.0, whose tasks have all completed, from p
16/06/02 21:34:18 INFO DAGScheduler: Job 0 finished: collect at /home/profesor/querySparkSQL.py:1
Row(_c0=59793)
16/06/02 21:34:18 INFO SparkContext: Invoking stop() from shutdown hook
16/06/02 21:34:18 INFO ContextHandler: stopped o.s.j.s.ServletContextHandler{/static/sql,null}
...
16/06/02 21:34:19 INFO MetricsSystemImpl: azure-file-system metrics system stopped.
16/06/02 21:34:19 INFO MetricsSystemImpl: azure-file-system metrics system shutdown complete.
profesor@hn0-ejempl:~$
```

Ejemplo: SparkSQL + Hive

```
hive> select count(*) from hivesampletable;  
Query ID = profesor_20160602214152_fa75b802-a9f0-4a39-8657-b176064094db  
Total jobs = 1  
Launching Job 1 out of 1
```

Status: Running (Executing on YARN cluster with App id application_1464884087522_0013)

```
-----  
VERTICES    STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED  
-----  
Map 1 ..... SUCCEEDED    1          1          0          0          0          0  
Reducer 2 ..... SUCCEEDED    1          1          0          0          0          0  
-----  
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 17.58 s  
-----
```

Status: DAG finished successfully in 17.59 seconds

METHOD	DURATION(ms)
parse	1,840
semanticAnalyze	2,805
TezBuildDag	935
TezSubmitToRunningDag	432
TotalPrepTime	8,408

VERTICES	TOTAL_TASKS	FAILED_ATTEMPTS	KILLED_TASKS	DURATION_SECONDS	CPU_TIME_MILLIS	GC_TIME_MILLIS	INPUT_RECORDS
Map 1	1	0	0	4.86	6,340	0	59,793
Reducer 2	1	0	0	2.12	2,650	0	1

```
OK  
59793  
Time taken: 16.896 seconds, Fetched: 1 row(s)  
hive>
```

SparkSQL y Hive
coinciden!!!

SQL + RDDs

```
from pyspark import SparkConf, SparkContext, SQLContext
from pyspark.sql.types import StringType

# Import Spark SQL
from pyspark.sql import HiveContext, Row
# Or if you can't include the hive requirements
from pyspark.sql import SQLContext, Row

conf = SparkConf().setAppName("GenerateTraces")
sc = SparkContext(conf = conf)
hiveCtx = HiveContext(sc)
query = hiveCtx.sql("SELECT clientid, devicemake, country FROM hivesampletable")
#for x in query.collect():
#    print x
keyValue = query.map(lambda x : (x.clientid,x.devicemake + ':' + x.country ))
groupByKV = keyValue.groupByKey().mapValues(list)
for x in groupByKV.collect():
    print x
```

SQL + RDDs

```
'Apple:United States', u'Apple:United States', u'Apple:United States', u'Apple:United States', u'Apple:U
(u'98189', [u'LG:United States'])
(u'21204', [u'Apple:United States', u'Apple:United States'])
(u'134260', [u'Apple:United States', u'Apple:United States', u'Apple:United States', u'Apple:United Stat
(u'45569', [u'LG:United States'])
(u'94389', [u'RIM:United States', u'RIM:United States', u'RIM:United States', u'RIM:United States'])
(u'4958', [u'Apple:United States'])
(u'43602', [u'Samsung:United States', u'Samsung:United States'])
(u'4952', [u'Apple:United States', u'Apple:United States'])
(u'43608', [u'RIM:United States'])
(u'62580', [u'Apple:United States', u'Apple:United States', u'Apple:United States'])
(u'17047', [u'RIM:United States', u'RIM:United States'])
(u'3298', [u'Samsung:United States'])
(u'96873', [u'Apple:United States', u'Apple:United States'])
(u'13207', [u'LG:United States', u'LG:United States', u'LG:United States', u'LG:United States'])
(u'13209', [u'Samsung:United States', u'Samsung:United States', u'Samsung:United States', u'Samsung:Unit
United States'])
(u'36383', [u'Apple:United States'])
(u'90093', [u'LG:United States', u'LG:United States', u'LG:United States', u'LG:United States', u'LG:Uni
(u'90053', [u'Samsung:United States', u'Samsung:United States', u'Samsung:United States', u'Samsung:Unit
(u'90763', [u'Apple:United States', u'Apple:United States', u'Apple:United States', u'Apple:United State
(u'136284', [u'Samsung:United States', u'Samsung:United States', u'Samsung:United States', u'Samsung:Uni
```

ELEGIMOS TODO

Tez versus Spark

Tez vs Spark

Apples vs. Oranges

<https://www.xplenty.com/blog/2015/01/apache-spark-vs-tez-comparison/>

This is how each framework brands itself:

“Apache Spark is a fast and general engine for large-scale data processing.” ([source](#))

“The Apache Tez project is aimed at building an application framework which allows for a complex directed-acyclic-graph of tasks for processing data. It is currently built atop Apache Hadoop YARN.” ([source](#))

Considering the fact that Spark also uses directed-acyclic-graphs, don't they sound a bit similar? Maybe. Nonetheless, in an [interview with Shaun Connolly](#), Hortonworks product strategy vice president, he differentiates between the two by saying that Spark is a general purpose engine with APIs for mainstream developers, while Tez is a framework for purpose-built tools such as Hive and Pig.

Nuestro benchmark

- Eventos:

ID_123,TIME_001,ANTENNA_A

ID_235,TIME_002,ANTENNA_A

ID_123,TIME_003,ANTENNA_B

ID_654,TIME_005,ANTENNA_D

...

ID_ABC,TIME_XYZ,ANTENNA_F

- Trazas de movilidad:

{ID_123,((TIME_001,ANTENNA_A),(TIME_003,ANTENNA_B))}

{ID_235,((TIME_002,ANTENNA_A))}

{ID_654,((TIME_005,ANTENNA_D))}

...

ELEGIMOS TODO_

Nuestro benchmark

- Tres implementaciones
 - Pig + Tez
 - Python Spark (interpretado)
 - Scala Spark (compilado)
- ¿Qué implementación permite calcular trazas de la manera más rápida?
- Intuición:
 - Código compilado debería ser más rápido

Nuestro benchmark

Datos: 10GBytes

	Pig+Tez	PySpark	Scala
Tiempo [mm:ss]	27:22	31:56	28:33
Líneas de código	24	55	35 + 7 (makefile)
Tuning [Horas]	0	48	48

Nuestro benchmark

Datos: 10GBytes

	Pig+Tez	PySpark	Scala
Tiempo [mm:ss]	27:22	31:56	28:33
Líneas de código	24	55	35 + 7 (makefile)
Tuning [Horas]	0	48	48

¿Tez o Spark?
Depende...

Gracias!!